

King Saud University
College of Computer and Information Sciences
Software Engineering department


SW 485: Selected topics in software engineering

Stroke Prediction Analysis

Project Report

Phase #1

Group:4



| Name | ID |
|------------------|-----------|
| Jouri Alanazi | 441200780 |
| Renad Alsubaie | 441200868 |
| Monerah Alsubaie | 441200880 |

Supervised BY: Dr.Manal Bin Khonain

Contents

| | |
|------------------------------------------|---|
| 1. Introduction: | 3 |
| Motivation:..... | 3 |
| Objective and Expected Results: | 3 |
| 2. Machine learning Tasks: | 4 |
| 3. Data: | 5 |
| 1.Data Kind:..... | 5 |
| 2.Data Source:..... | 5 |
| 3.Data Exploration: | 5 |
| Number of observations: | 5 |
| Number of variables and data type: | 5 |
| 4. Data preprocessing..... | 6 |

1. Introduction:

Machine learning is a subfield of artificial intelligence that involves the use of algorithms and statistical models to analyze and interpret complex data. In recent years, machine learning has become increasingly popular in the healthcare industry, particularly in the field of stroke prediction analysis.

Motivation:

A stroke is a serious medical condition that occurs when blood flow to the brain is interrupted, either due to a blockage in the blood vessels or bleeding in the brain. Early detection of stroke is crucial for timely treatment, which can help prevent serious complications and even death.

Machine learning algorithms can be used to analyze large datasets of patient information to identify patterns and risk factors that may contribute to stroke. By analyzing factors such as age, sex, medical history, and lifestyle habits, machine learning models can predict a patient's risk of stroke and help healthcare providers develop personalized treatment plans.

Objective and Expected Results:

The objective of using machine learning in stroke prediction analysis is to improve the accuracy and efficiency of stroke risk assessment. By analyzing large amounts of patient data using machine learning algorithms, healthcare providers can identify patterns and risk factors that may not be immediately apparent through traditional methods.

The expected results of using machine learning in stroke prediction analysis include:

1. Improved accuracy: Machine learning algorithms can identify subtle patterns in patient data that may not be immediately apparent through traditional methods. This can lead to more accurate stroke risk assessments and better treatment decisions.
2. Early detection: Machine learning models can be trained to detect early warning signs of stroke, allowing for timely interventions that can prevent serious complications.
3. Personalized treatment: By analyzing patient data, machine learning algorithms can help healthcare providers develop personalized treatment plans based on individual risk factors and medical history.
4. Reduced healthcare costs: Early detection and personalized treatment can help reduce healthcare costs by preventing costly complications and hospitalizations.

Overall, the use of machine learning in stroke prediction analysis has the potential to improve patient outcomes and advance medical research in the field of stroke prevention and treatment.

2. Machine learning Tasks:

In stroke prediction analysis, machine learning can be used for several tasks, including:

1. **Feature selection:** Machine learning algorithms can identify the most relevant features or risk factors that contribute to stroke prediction. This involves selecting the most important variables from a large dataset of patient information that can be used to develop a predictive model.
2. **Classification:** Machine learning algorithms can be used to classify patients into different risk categories based on their individual risk factors. This allows healthcare providers to identify patients who are at a higher risk of stroke and develop personalized prevention plans accordingly.
3. **Regression:** Machine learning algorithms can also be used to predict the probability of stroke occurrence or the time to stroke occurrence for individual patients. This can help healthcare providers develop more accurate and personalized treatment plans.
4. **Clustering:** Machine learning algorithms can group patients with similar risk profiles together, which can help healthcare providers identify common risk factors and develop more effective prevention strategies.
5. **Anomaly detection:** Machine learning algorithms can be used to identify patients with unusual or unexpected risk factors that may not be immediately apparent through traditional methods. This can help healthcare providers identify patients who may be at a higher risk of stroke and develop personalized prevention plans accordingly.

Overall, machine learning can be a powerful tool for stroke prediction analysis, allowing healthcare providers to develop more accurate and personalized prevention and treatment plans for individual patients.

3. Data:

1.Data Kind:

contains information about patients who may or may not have experienced a stroke. The dataset includes demographic information such as age, gender, and smoking status, as well as medical history such as hypertension, heart disease, and glucose levels. The target variable is whether or not the patient has had a stroke.

The dataset contains 5110 observations and 12 variables, This dataset can be used for modeling and predicting the likelihood of a patient having a stroke based on their demographic and medical information. The dataset is useful for researchers and healthcare professionals interested in developing stroke prevention and treatment strategies.

2.Data Source:

The Stroke Prediction Dataset available on Kaggle was collected and made available .The data was collected from patients who visited a healthcare facility in a rural area of India. The information was gathered through medical examinations and interviews with the patients.

We got the dataset from: <https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>

3.Data Exploration:

Number of observations:

The Stroke Prediction Dataset contains 5110 observations or rows. Each row represents a patient in the dataset.

Number of variables and data type:

The Stroke Prediction Dataset contains 12 variables or columns. Here is the list of variables and their data types:

1. id: Numeric - unique identifier for each patient
2. gender: Categorical - male or female
3. age: Numeric - age of the patient in years
4. hypertension: Categorical - 0 if the patient doesn't have hypertension, 1 if the patient has hypertension
5. heart disease: Categorical - 0 if the patient doesn't have any heart diseases, 1 if the patient has a heart disease
6. ever married: Categorical - "No" or "Yes"
7. work type: Categorical - "children", "Govt jov", "Never worked", "Private" or "Self-employed"
8. Residence type: Categorical - "Rural" or "Urban"
9. avg glucose level: Numeric - average glucose level in the patient's blood
10. bmi: Numeric - body mass index of the patient
11. smoking status: Categorical - "formerly smoked", "never smoked", "smokes" or "Unknown"
12. stroke: Categorical - 1 if the patient had a stroke, 0 if not

This Figure displays the data type for each variable in our dataset method:

```
In [12]: #display datatypes for each column
data.dtypes
```

```
Out[12]: gender      object
age      float64
hypertension  int64
heart_disease  int64
ever_married  object
work_type  object
Residence_type  object
avg_glucose_level  float64
bmi      float64
smoking_status  object
stroke    int64
dtype: object
```

Figure 1. Variables and their data types in the dataset

Sample of raw dataset:

To display a sample from the dataset, we also used the. head () [1] method. The first five rows of the dataset were displayed by the function.

```
In [50]: #Sample of raw dataset|
#Display The First Five Information in Stroke Prediction Dataset
data.head()
```

```
Out[50]:
```

| | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | stroke |
|-------|--------|------|--------------|---------------|--------------|---------------|----------------|-------------------|------|-----------------|--------|
| id | | | | | | | | | | | |
| 9046 | Male | 67.0 | 0 | 1 | Yes | Private | Urban | 228.69 | 36.6 | formerly smoked | 1 |
| 51676 | Female | 61.0 | 0 | 0 | Yes | Self-employed | Rural | 202.21 | NaN | never smoked | 1 |
| 31112 | Male | 80.0 | 0 | 1 | Yes | Private | Rural | 105.92 | 32.5 | never smoked | 1 |
| 60182 | Female | 49.0 | 0 | 0 | Yes | Private | Urban | 171.23 | 34.4 | smokes | 1 |
| 1665 | Female | 79.0 | 1 | 0 | Yes | Self-employed | Rural | 174.12 | 24.0 | never smoked | 1 |

Figure 2.Raw Sample from a dataset

Displaying Graphs in Dataset:

```
In [115]: #Corrleation and we describe it with grapg and if it's became more brighter then there is a relationship between the  
plt.figure(figsize=(20, 10))  
sns.heatmap(data.corr(numeric_only = True) , annot=True , cmap="Blues" , linecolor="White")
```

Out [115]: <Axes: >



Figure 3. Correlation for the dataset

Variables Distribution

Distribution plot

This distribution plot depicts the overall distribution and shows the impact of age and avg_glucose_level and BMI.

```
In [52]: #Variables Distribution - Figure 1
#Histogram Density
column_list = ['age', 'avg_glucose_level', 'bmi']

fig, ax = plt.subplots(1, 3, figsize=(12,6))

for i, col in enumerate(column_list):
    sns.histplot(data=data, x=col, stat='density', hue='stroke', multiple='stack', ax=ax[i])
    sns.kdeplot(data=data, x=col, linewidth=3, hue='stroke', fill=False, multiple='stack', ax=ax[i])

plt.tight_layout()
plt.show()
```

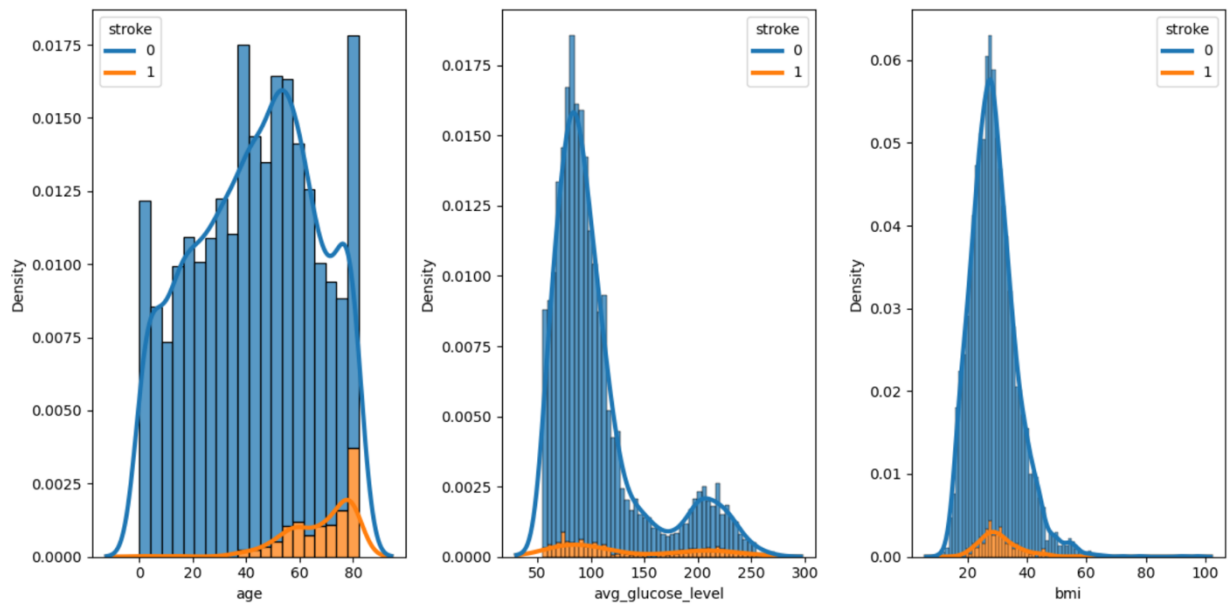


Figure 4. Distribution plot

Scatter plot

Scatter plot is usually used to see the correlation between attributes we want to know the correlation between patients age and average glucose level.

```
In [51]: #Variables Distribution – Figure 2
#Scatter Plot
plt.figure(figsize=(12,6))
sns.jointplot(data=data, x='age', y='avg_glucose_level')
plt.show()
```

<Figure size 1200x600 with 0 Axes>




Figure 5.Joint plot

Pie Chart

Another way to see the portion of categories is using a pie chart. But, we must be very careful when using this chart type.

```
In [82]: #Variables Distribution - Figure 4
# Create a pie chart
plt.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%', startangle=90)
plt.axis('equal')
plt.title('Stroke vs No Stroke')
plt.show()
```




Figure 6. Pie Chart

Bar Plot

The Bar chart compares between the two type of gender (Female & Male).

```
In [53]: #Variables Distribution - Figure 3
#Bar Plot
data['gender'].value_counts().plot(kind='bar')
print(data['gender'].value_counts())
plt.title('Distribution of gender', fontsize=15)
```

```
Female    2994
Male      2115
Other         1
Name: gender, dtype: int64
```

```
Out[53]: Text(0.5, 1.0, 'Distribution of gender')
```




Figure 7.Bar Plot

Missing values(null):

To Detect any missing values we use. `isnull()[1]` function from our dataset:

```
In [83]: #Missing Values
data.isnull()
```

```
Out[83]:
```

| | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | stroke |
|-------|--------|-------|--------------|---------------|--------------|-----------|----------------|-------------------|-------|----------------|--------|
| id | | | | | | | | | | | |
| 9046 | False | False | False | False | False | False | False | False | False | False | False |
| 51676 | False | False | False | False | False | False | False | False | True | False | False |
| 31112 | False | False | False | False | False | False | False | False | False | False | False |
| 60182 | False | False | False | False | False | False | False | False | False | False | False |
| 1665 | False | False | False | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 18234 | False | False | False | False | False | False | False | False | True | False | False |
| 44873 | False | False | False | False | False | False | False | False | False | False | False |
| 19723 | False | False | False | False | False | False | False | False | False | False | False |
| 37544 | False | False | False | False | False | False | False | False | False | False | False |
| 44679 | False | False | False | False | False | False | False | False | False | False | False |

5110 rows x 11 columns

Figure 8. Display Null Values

Statistical Summaries:

In this table we calculate mean, std, min, max for the numeric columns we will use. `describe()[1]` function from our dataset:

```
In [13]: #Statistical Summaries
#compute the staticial summaries for all variables
data.describe()
```

```
Out[13]:
```

| | age | hypertension | heart_disease | avg_glucose_level | bmi | stroke |
|-------|-------------|--------------|---------------|-------------------|-------------|-------------|
| count | 5110.000000 | 5110.000000 | 5110.000000 | 5110.000000 | 4909.000000 | 5110.000000 |
| mean | 43.226614 | 0.097456 | 0.054012 | 106.147677 | 28.893237 | 0.048728 |
| std | 22.612647 | 0.296607 | 0.226063 | 45.283560 | 7.854067 | 0.215320 |
| min | 0.080000 | 0.000000 | 0.000000 | 55.120000 | 10.300000 | 0.000000 |
| 25% | 25.000000 | 0.000000 | 0.000000 | 77.245000 | 23.500000 | 0.000000 |
| 50% | 45.000000 | 0.000000 | 0.000000 | 91.885000 | 28.100000 | 0.000000 |
| 75% | 61.000000 | 0.000000 | 0.000000 | 114.090000 | 33.100000 | 0.000000 |
| max | 82.000000 | 1.000000 | 1.000000 | 271.740000 | 97.600000 | 1.000000 |

Figure 9. Display Statistical Summaries

4. Data preprocessing

The Stroke Prediction Dataset contains medical and demographic data for over 5000 patients, with the goal of predicting whether a patient is at risk of having a stroke based on their medical and demographic information.

Before we can model the data, it's important to preprocess it to ensure its quality and suitability for analysis. The preprocessing steps for this dataset may include:

1. Handling missing values: The dataset contains some missing values represented by "N/A" or "Unknown". These missing values can be handled by either imputing the missing values using statistical methods or dropping the rows with missing values if the missing values are too many.
2. Encoding categorical variables: The categorical variables in the dataset such as gender, ever_married, work_type, Residence_type, and smoking_status need to be encoded as numeric values to be used in modeling. This can be done using methods such as one-hot encoding or label encoding.
3. Scaling numeric variables: The numeric variables in the dataset such as age, avg_glucose_level, and bmi may need to be scaled to have a similar range of values. This can be done using methods such as standardization or normalization.
4. Handling imbalanced data: The dataset is imbalanced with a majority of patients not having a stroke. This can lead to biased models. Techniques such as oversampling the minority class or undersampling the majority class can be used to balance the data.
5. Feature selection: Not all variables in the dataset may be relevant for modeling. Feature selection techniques can be used to select the most important variables for predicting stroke.

By preprocessing the data, we can improve the quality of the data and make it suitable for modeling. We can then use various machine learning algorithms to predict stroke risk in patients and help healthcare providers make informed decisions.

And we check if there are any duplicate values in our dataset and we did not find and below the snapshot.

```
In [87]: # Check for duplicates
duplicates = data.duplicated()
num_duplicates = duplicates.sum()
print('Number of duplicates: {}'.format(num_duplicates))

Number of duplicates: 0
```

Figure 10. Snapshot to check if there is duplicate data

Also, we make discretization for our dataset between age and BMI.

```
In [90]: #Discretization Code
# Define the bins for age and bmi
age_bins = [0, 18, 35, 50, 65, 100]
bmi_bins = [0, 18.5, 25, 30, 35, 100]

# Discretize the age and bmi columns using the cut function
data['age_group'] = pd.cut(data['age'], bins=age_bins, labels=['<18', '18-35', '35-50', '50-65', '65+'])

data['bmi_group'] = pd.cut(data['bmi'], bins=bmi_bins, labels=['Underweight', 'Normal', 'Overweight', 'Obese', 'Extremely Obese'])

# Print the first few rows of the discretized data
print(data[['age', 'age_group', 'bmi', 'bmi_group']].head())
```

| | age | age_group | bmi | bmi_group |
|-------|------|-----------|------|-----------------|
| id | | | | |
| 9046 | 67.0 | 65+ | 36.6 | Extremely Obese |
| 51676 | 61.0 | 50-65 | NaN | NaN |
| 31112 | 80.0 | 65+ | 32.5 | Obese |
| 60182 | 49.0 | 35-50 | 34.4 | Obese |
| 1665 | 79.0 | 65+ | 24.0 | Normal |

Figure 11.Snapshot of Discretization Method

And we also make clean for our code, so we try to delete a row from our dataset and we display our dataset in several Snapshots by displaying our dataset before and after clean and also during using the method for deleting the row.

```
In [5]: #Data Preprocessing
#Stroke prediction dataset before delete the row contain 'smokes'
data.head()
```

```
Out[5]:
```

| | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | stroke |
|-------|--------|------|--------------|---------------|--------------|---------------|----------------|-------------------|------|-----------------|--------|
| id | | | | | | | | | | | |
| 9046 | Male | 67.0 | 0 | 1 | Yes | Private | Urban | 228.69 | 36.6 | formerly smoked | 1 |
| 51676 | Female | 61.0 | 0 | 0 | Yes | Self-employed | Rural | 202.21 | NaN | never smoked | 1 |
| 31112 | Male | 80.0 | 0 | 1 | Yes | Private | Rural | 105.92 | 32.5 | never smoked | 1 |
| 60182 | Female | 49.0 | 0 | 0 | Yes | Private | Urban | 171.23 | 34.4 | smokes | 1 |
| 1665 | Female | 79.0 | 1 | 0 | Yes | Self-employed | Rural | 174.12 | 24.0 | never smoked | 1 |

Figure 12.Snapshot for the dataset before delete

```
In [6]: #Data Preprocessing
#code for delete the row contain 'smokes'
data = data[(data['smoking_status'] != 'smokes')]
```

Figure 13.Snapshot for the dataset during deleting the row

```
In [7]: #Data Preprocessing
#display the dataset after delete the row contain 'smokes'
data.head()
```

```
Out[7]:
```

| | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | stroke |
|-------|--------|------|--------------|---------------|--------------|---------------|----------------|-------------------|------|-----------------|--------|
| id | | | | | | | | | | | |
| 9046 | Male | 67.0 | 0 | 1 | Yes | Private | Urban | 228.69 | 36.6 | formerly smoked | 1 |
| 51676 | Female | 61.0 | 0 | 0 | Yes | Self-employed | Rural | 202.21 | NaN | never smoked | 1 |
| 31112 | Male | 80.0 | 0 | 1 | Yes | Private | Rural | 105.92 | 32.5 | never smoked | 1 |
| 1665 | Female | 79.0 | 1 | 0 | Yes | Self-employed | Rural | 174.12 | 24.0 | never smoked | 1 |
| 56669 | Male | 81.0 | 0 | 0 | Yes | Private | Urban | 186.21 | 29.0 | formerly smoked | 1 |

Figure 14.Snapshot for the dataset after clean

6. References:

- [1] <https://www.geeksforgeeks.org/data-preprocessing-in-data-mining/>
- [2] Ganpathi, A. (2021) *Data preprocessing and Exploratory Data Analysis for machine learning.*, Medium. Analytics Vidhya. Available at: <https://medium.com/analytics-vidhya/data-preprocessing-and-exploratory-data-analysis-for-machine-learning-75b8a6468b72> (Accessed: April 26, 2023).
- [3] Fedesoriano (2021) *Stroke prediction dataset*, Kaggle. Available at: <https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset/code?datasetId=1120859&searchQuery=basic> (Accessed: April 26, 2023).